

# Zhiheng Zhang

m15601654187@163.com | zhihezhang@student.unimelb.edu.au  
github.com/zhiheng-zhang-Mera

Melbourne, Australia

## Research Profile

Computer Science researcher and research engineer interested in **trustworthy scientific agents, computational health workflows, and evidence-based multi-agent reasoning**. My current work on multi-agent orchestration motivates a broader question: how autonomous systems can synthesize heterogeneous evidence, revise hypotheses, interact with analytical tools, and make their conclusions inspectable.

## Education

### University of Melbourne

*Master's degree in Computer Science*

Melbourne, Australia

*Expected 2026*

### University of British Columbia

*Bachelor of Science in Computer Science*

Vancouver, Canada

*Completed Summer 2024*

## Selected Research and Engineering Projects

### Codex Boss — *Multi-agent orchestration and evidence-based acceptance platform*

2026–present

github.com/zhiheng-zhang-Mera/Codex-Boss

- Building a platform for long-running AI workflows in which multiple agents can plan, execute, critique, and adjudicate work while preserving durable task state.
- Developed platform foundations spanning durable state/events, capability boundaries, plugin isolation, knowledge/data lifecycle, dogfooding, and acceptance/verification.
- Investigating how autonomous systems can distinguish genuine task completion from superficially plausible but weak evidence.
- Current semantic-acceptance work deliberately retains unresolved failure cases: a vacuous case is rejected, while a meaningful case still exposes a false negative in the evidence reader rather than being reported as solved.
- Current regression evidence on the relevant branch includes **2,508 unit tests across 215 files**, plus Electron, renderer, and test typechecking.

### Computational Health Prototyping — *Scientific-workflow and simulation concepts*

- Prototyping computational-health and simulation ideas that represent physiological context, interventions, mechanisms, time-dependent effects, and uncertainty as structured inputs and outputs.
- Using these prototypes to explore how agent systems can support evidence synthesis, hypothesis revision, simulation/data-analysis workflows, and inspectable scientific reasoning.

### DS-Hns — *Long-horizon autonomous software-work runtime*

2026–present

github.com/zhiheng-zhang-Mera/DS-Hns

- Building an execution layer for unattended software work around the official DeepSeek Harness interface, with optional capabilities isolated as extensions.
- Implemented ordered/scheduled task queues, hardware-adaptive local concurrency, unified task lifecycle and terminal events, task history, and terminal notifications.
- Designed extension-failure isolation, restart/recovery-oriented execution, and an idempotent Windows installer with explicit optional-plugin flow.
- Using the system as a testbed for **reliable coding agents, persistent execution, recovery, scheduling, and completion verification**.

### Privacy Lens — *Trustworthy-software / privacy project*

- Explored privacy-risk interpretation and user-facing warning mechanisms with attention to bounded interpretation, auditability, evidence boundaries, and trustworthy software behavior.
- Supports broader interests in trustworthy AI, software governance, and high-stakes automated systems.

## Research Interests

Scientific agents; computational health and biomedical AI; evidence synthesis; multi-agent reasoning and critique; tool-using AI for scientific workflows; trustworthy evaluation of long-horizon research agents.

## Technical and Research Strengths

- End-to-end research engineering: problem formulation, system design, iterative implementation, testing, debugging, and empirical evaluation.

- Multi-module software design with explicit persistence, lifecycle, extensibility, isolation, recovery, and acceptance criteria.
- Git/GitHub-based experimental development, branch-based iteration, regression testing, and reproducible project documentation.
- Practical multi-AI workflow design: using multiple AI systems as coding/research collaborators while independently reviewing plans, evidence, failure modes, and completion claims.

## Current Research Direction

---

I want to study **trustworthy scientific agents for computational health**: systems that can integrate heterogeneous evidence, propose and revise hypotheses or experiment plans, interact with simulation/data-analysis tools, and expose inspectable evidence for their conclusions instead of merely generating plausible narratives.