

Zhiheng Zhang

zhiheng0mera@gmail.com | zhihezhang@student.unimelb.edu.au
github.com/zhiheng-zhang-Mera

Melbourne, Australia

Research Profile

Computer Science research engineer focused on reliable LLM agents, multi-agent evaluation, long-horizon execution and human-controlled agent systems.

Education

University of Melbourne

Master's degree in Computer Science

Melbourne, Australia

Expected 2026

University of British Columbia

Bachelor of Science in Computer Science

Canada

Completed Summer 2024

Selected Research and Engineering Projects

Codex Boss — *Agent orchestration and evidence-based acceptance*

2026–present

github.com/zhiheng-zhang-Mera/Codex-Boss

- Developed a multi-agent workflow platform for planning, execution, critique and adjudication with durable state and evidence-aware acceptance.
- Treats unresolved failure cases and incomplete acceptance conditions as explicit evidence rather than reporting them as solved.
- Supports trustworthy agent orchestration, completion claims, auditability and human oversight.

Utopia / Digital City — *Multi-device personal-computing and research platform*

2026–present

github.com/zhiheng-zhang-Mera/Utopia

- Built Android and Web control surfaces around a platform-neutral City protocol with authenticated device paths, explicit confirmation and truthful degraded-state handling.
- Current accepted product evidence includes a three-end mesh across two real Windows workers plus Android, with strict target-device routing and cross-host formal review.
- Research-strengthening stages for experiment registry, provenance/metrics, repeated scenarios, fault injection and replay/ablation have accepted evidence; artifact export remained under independent review at the 2026-10-06 snapshot.
- Planned next direction: a **Personal Compute Fabric** for heterogeneous-device telemetry, placement/offloading, recovery and explainable cross-device execution.

DS-Hns — *Long-horizon autonomous software-work runtime*

2026–present

github.com/zhiheng-zhang-Mera/DS-Hns

- Building long-horizon repository execution with durable task state, restart/recovery, hardware-aware concurrency and explicit terminal-state evidence.
- Uses failure isolation and recovery-oriented execution to study when a software agent has genuinely completed work rather than merely survived.
- Supports coding-agent reliability, autonomous maintenance and empirical software-engineering evaluation.

Research Interests

Reliable LLM agents; multi-agent evaluation; long-horizon agents; self-evolving agents; human-agent interaction; reproducible agent systems.

Technical and Research Strengths

- End-to-end research engineering across system design, implementation, testing, controlled experiments and evidence-driven iteration.
- Multi-device and long-running software systems with explicit lifecycle, recovery, routing, observability and acceptance boundaries.
- Git/GitHub-based reproducible development, regression testing, cross-host review and structured preservation of failure evidence.

Current Research Direction

I want to study agent systems that improve over long horizons without losing evidence, user control or the ability to reproduce failures across tools and devices.